

```
# ===== 0.1 Check for GPU
!nvidia-smi

import tensorflow as tf
from tensorflow.python.client import device_lib

print("Tensorflow Version:", tf.__version__)
print("Is Build with CUDA:", tf.test.is_built_with_cuda())
print("Is GPU available:", tf.test.is_gpu_available(cuda_only=False, min_cuda_compute_capability=None))
print(device_lib.list_local_devices())
```



Sat Feb 29 23:40:23 2020

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI		440.48.02		Driver Version: 418.67			CUDA Version: 10.1		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU	Name	Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr.	ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage		GPU-Util	Compute	M.	
=====+=====+=====+=====+=====+=====+=====+=====+=====+=====									
0	Tesla	P100-PCIE...	Off	00000000:00:04.0	Off			0	
N/A	33C	P0	26W / 250W	0MiB / 16280MiB		0%		Default	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

Processes:					GPU Memory
GPU	PID	Type	Process name		Usage
=====					
No running processes found					

The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x.

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %tensorflow_vers

Tensorflow Version: 1.15.0

Is Build with CUDA: True

Is GPU available: True

[name: "/device:CPU:0"

device_type: "CPU"

memory_limit: 268435456

locality {

}

incarnation: 11877076774216315395

, name: "/device:XLA_CPU:0"

device_type: "XLA_CPU"

memory_limit: 17179869184

locality {

}

incarnation: 15791819063071657308

physical_device_desc: "device: XLA_CPU device"

, name: "/device:XLA_GPU:0"

device_type: "XLA_GPU"

memory_limit: 17179869184

locality {

}

incarnation: 6432898729218651728

physical_device_desc: "device: XLA_GPU device"

, name: "/device:GPU:0"

device_type: "GPU"

memory_limit: 15956161332

locality {

bus_id: 1

links {

}

}

incarnation: 16446644522011177398

physical_device_desc: "device: 0, name: Tesla P100-PCIE-16GB, pci bus id: 0000:00:04.0, compute ca

]

===== 0.2 Utils and Consts

import time

import psutil

import pandas as pd

```

import pandas as pd
import numpy as np
import cv2
import base64
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.applications import vgg16, vgg19, resnet50
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D, GlobalAveragePooling2D
from tensorflow.keras.layers import BatchNormalization, Activation, Input, AveragePooling2D
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, LearningRateScheduler, ReduceLROnPlateau
from tensorflow.keras.regularizers import l2

```

```

PATH_ROOT = 'drive/My Drive/Colab Notebooks'
PATH_FILE = PATH_ROOT + '/Datasets/ChartImages/charts_1_day_50_periods.csv'
MODEL_NAME_RESNET = PATH_ROOT + '/Models/ChartImages/model_resnet_1_day_50_periods.h5'
MAX_RGB = 255
IMAGE_DIMENSION = 224
IMAGE_INPUT_SHAPE = (IMAGE_DIMENSION, IMAGE_DIMENSION, 3)
LABEL_CLASS = { 'Down': 0, 'Up': 1 }
NUM_CLASSES = 2
NUM_BATCH_SIZE_RESNET = 32
NUM_DEPTH = 29
NUM_EPOCHS = 100

```

```
# Utils
```

```

start_time = time.time()
def watch_restart():
    global start_time
    start_time = time.time()
def watch_print(title):
    global start_time
    print(title, round(time.time() - start_time, 4), 'seconds')
def memory_print():
    memory = dict(psutil.virtual_memory()._asdict())
    print("Memory Capacity", memory['total'] >> 30, "GB")
    print("Memory Left", memory['free'] >> 30, "GB")
    print("Memory Used", memory['used'] >> 30, "GB")
    print("Memory Used:", memory['percent'], "percent")
def lr_schedule(epochs):
    lr = 1e-3
    if epochs > 180:
        lr *= 0.5e-3
    elif epochs > 160:
        lr *= 1e-3
    elif epochs > 120:
        lr *= 1e-2
    elif epochs > 80:
        lr *= 1e-1
    return lr

```

```
# Callbacks for model saving and stopping.
```

```

# Training should be stopped when val_acc (validation accuracy) clearly stops increasing to prevent overfitting
checkpoint_resnet = ModelCheckpoint(filepath=MODEL_NAME_RESNET, monitor='val_acc', verbose=1, save_best_only=True)
lr_scheduler = LearningRateScheduler(lr_schedule)

```

```
lr_reducer = ReduceLROnPlateau(factor=np.sqrt(0.1), cooldown=0, patience=5, min_lr=0.5e-6)
early_stopping = EarlyStopping(monitor='val_acc', patience=10, verbose=1)
```

```
watch_restart()
```

```
# ===== 1.0 Get Image Data from CSV
```

```
df = pd.read_csv(PATH_FILE)
print(df.head())
```

```
memory_print()
```

```
watch_print('Get Data')
```

```

└─ Id Symbol ... Y_Prediction X_Image
0 1 M ... 1 iVBORw0KGgoAAAANSUhEUgAAOAAAADgCAYAAAAaLWrhAA...
1 2 M ... 1 iVBORw0KGgoAAAANSUhEUgAAOAAAADgCAYAAAAaLWrhAA...
2 3 M ... 1 iVBORw0KGgoAAAANSUhEUgAAOAAAADgCAYAAAAaLWrhAA...
3 4 M ... 1 iVBORw0KGgoAAAANSUhEUgAAOAAAADgCAYAAAAaLWrhAA...
4 5 M ... 1 iVBORw0KGgoAAAANSUhEUgAAOAAAADgCAYAAAAaLWrhAA...

```

```

[5 rows x 5 columns]
Memory Capacity 25 GB
Memory Left 21 GB
Memory Used 1 GB
Memory Used: 4.8 percent
Get Data 2.6114 seconds

```

```
watch_restart()
```

```
# ===== 2.0 Prepare Data
```

```
# Set up X and y
```

```
items = []
```

```
for index, row in df.iterrows():
```

```
    # Convert from base64 string to byte array
```

```
    item_byte_array = base64.b64decode(df['X_Image'][index])
```

```
    # Convert byte array to numpy array for OpenCv usage
```

```
    item_np = np.frombuffer(item_byte_array, dtype=np.uint8)
```

```
    # Convert numpy array to OpenCv image
```

```
    item_image = cv2.imdecode(item_np, flags=1)
```

```
    items.append(item_image)
```

```
    # if index < 1:
```

```
        # plt.imshow(cv2.cvtColor(item_image, cv2.COLOR_BGR2RGB));
```

```
        # plt.show()
```

```
X = np.array(items)
```

```
y = to_categorical(df[['Y_Prediction']].values)
```

```
df = None # Clear RAM
```

```
print('y type', type(y))
```

```
print('y shape', y.shape)
```

```
print('X type', type(X))
```

```
print('X shape', X.shape)
```

```
# Normalize input data.
```

```

# Neural Networks work best when input data are between 0 and 1 (Instead of 0 to 255).
X = X / MAX_RGB

# Split Train and Test
def split(X, y, proportion):
    ratio = int(X.shape[0]/proportion)
    X_train = X[ratio:,:]
    X_test = X[:ratio,:]
    y_train = y[ratio:,:]
    y_test = y[:ratio,:]
    return X_train, X_test, y_train, y_test
# X_train, X_test, y_train, y_test = split(X, y, 4) # Uses less RAM
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42) # Uses a lot of RAM
X = y = None # Clear RAM

print("X_train Shape", X_train.shape)
print("y_train Shape", y_train.shape)
print("X_test Shape", X_test.shape)
print("y_test Shape", y_test.shape)

memory_print()
watch_print('Prepare Data')

y type <class 'numpy.ndarray'>
y shape (9738, 2)
X type <class 'numpy.ndarray'>
X shape (9738, 224, 224, 3)
X_train Shape (7303, 224, 224, 3)
y_train Shape (7303, 2)
X_test Shape (2435, 224, 224, 3)
y_test Shape (2435, 2)
Memory Capacity 25 GB
Memory Left 0 GB
Memory Used 24 GB
Memory Used: 52.9 percent
Prepare Data 19.2904 seconds

watch_restart()

# ===== 3. Create Model Resnet
def resnet_layer(inputs,
                  num_filters=16,
                  kernel_size=3,
                  strides=1,
                  activation='relu',
                  batch_normalization=True,
                  conv_first=True):

    conv = Conv2D(num_filters,
                  kernel_size=kernel_size,
                  strides=strides,
                  padding='same',
                  kernel_initializer='he_normal',
                  kernel_regularizer=l2(1e-4))

    x = inputs
    if conv_first:
        x = conv(x)

```

```

    if batch_normalization:
        x = BatchNormalization()(x)
    if activation is not None:
        x = Activation(activation)(x)
else:
    if batch_normalization:
        x = BatchNormalization()(x)
    if activation is not None:
        x = Activation(activation)(x)
    x = conv(x)
return x

def get_model_resnet(input_shape, depth, num_classes):
    """ResNet

    Stacks of (1 x 1)-(3 x 3)-(1 x 1) BN-ReLU-Conv2D or also known as
    bottleneck layer
    First shortcut connection per layer is 1 x 1 Conv2D.
    Second and onwards shortcut connection is identity.
    At the beginning of each stage, the feature map size is halved (downsampled)
    by a convolutional layer with strides=2, while the number of filter maps is
    doubled. Within each stage, the layers have the same number filters and the
    same filter map sizes.
    Features maps sizes:
    conv1 : 32x32, 16
    stage 0: 32x32, 64
    stage 1: 16x16, 128
    stage 2: 8x8, 256

    """
    if (depth - 2) % 9 != 0:
        raise ValueError('depth should be 9n+2 (eg 56 or 110 in [b])')
    # Start model definition.
    num_filters_in = 16
    num_res_blocks = int((depth - 2) / 9)

    inputs = Input(shape=input_shape)
    # v2 performs Conv2D with BN-ReLU on input before splitting into 2 paths
    x = resnet_layer(inputs=inputs,
                     num_filters=num_filters_in,
                     conv_first=True)

    # Instantiate the stack of residual units
    for stage in range(3):
        for res_block in range(num_res_blocks):
            activation = 'relu'
            batch_normalization = True
            strides = 1
            if stage == 0:
                num_filters_out = num_filters_in * 4
                if res_block == 0: # first layer and first stage
                    activation = None
                    batch_normalization = False
            else:
                num_filters_out = num_filters_in * 2
                if res_block == 0: # first layer but not first stage
                    strides = 2 # downsample

```

```

# bottleneck residual unit
y = resnet_layer(inputs=x,
                  num_filters=num_filters_in,
                  kernel_size=1,
                  strides=strides,
                  activation=activation,
                  batch_normalization=batch_normalization,
                  conv_first=False)

y = resnet_layer(inputs=y,
                  num_filters=num_filters_in,
                  conv_first=False)

y = resnet_layer(inputs=y,
                  num_filters=num_filters_out,
                  kernel_size=1,
                  conv_first=False)

if res_block == 0:
    # linear projection residual shortcut connection to match
    # changed dims
    x = resnet_layer(inputs=x,
                      num_filters=num_filters_out,
                      kernel_size=1,
                      strides=strides,
                      activation=None,
                      batch_normalization=False)

x = tf.keras.layers.add([x, y])

num_filters_in = num_filters_out

# Add classifier on top.
# Has BN-ReLU before Pooling
x = BatchNormalization()(x)
x = Activation('relu')(x)
x = AveragePooling2D(pool_size=8)(x)
y = Flatten()(x)
outputs = Dense(num_classes,
                 activation='softmax',
                 kernel_initializer='he_normal')(y)

# Instantiate model.
model = Model(inputs=inputs, outputs=outputs)
return model

model_resnet = get_model_resnet(input_shape=IMAGE_INPUT_SHAPE,
                                depth=NUM_DEPTH,
                                num_classes=NUM_CLASSES)

model_resnet.compile(loss='categorical_crossentropy',
                    optimizer=Adam(lr=lr_schedule(NUM_EPOCHS)),
                    metrics=['accuracy'])

model_resnet.summary()

history_resnet = model_resnet.fit(X_train,
                                  y_train,
                                  batch_size=NUM_BATCH_SIZE_RESNET,
                                  epochs=NUM_EPOCHS,
                                  validation_data=(X_test, y_test),
                                  shuffle=True,
                                  callbacks=[checkpoint_resnet, lr_reducer, lr_scheduler, early_stopping])

```

```
watch_print('Create Model Resnet')
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/resource_instructions.py:166: instructions for updating:

If using Keras pass *_constraint arguments to layers.

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	
conv2d (Conv2D)	(None, 224, 224, 16)	448	input_1[0][0]
batch_normalization (BatchNormaliza	(None, 224, 224, 16)	64	conv2d[0][0]
activation (Activation)	(None, 224, 224, 16)	0	batch_normalization[0][0]
conv2d_1 (Conv2D)	(None, 224, 224, 16)	272	activation[0][0]
batch_normalization_1 (BatchNor	(None, 224, 224, 16)	64	conv2d_1[0][0]
activation_1 (Activation)	(None, 224, 224, 16)	0	batch_normalization_1[0][0]
conv2d_2 (Conv2D)	(None, 224, 224, 16)	2320	activation_1[0][0]
batch_normalization_2 (BatchNor	(None, 224, 224, 16)	64	conv2d_2[0][0]
activation_2 (Activation)	(None, 224, 224, 16)	0	batch_normalization_2[0][0]
conv2d_4 (Conv2D)	(None, 224, 224, 64)	1088	activation[0][0]
conv2d_3 (Conv2D)	(None, 224, 224, 64)	1088	activation_2[0][0]
add (Add)	(None, 224, 224, 64)	0	conv2d_4[0][0] conv2d_3[0][0]
batch_normalization_3 (BatchNor	(None, 224, 224, 64)	256	add[0][0]
activation_3 (Activation)	(None, 224, 224, 64)	0	batch_normalization_3[0][0]
conv2d_5 (Conv2D)	(None, 224, 224, 16)	1040	activation_3[0][0]
batch_normalization_4 (BatchNor	(None, 224, 224, 16)	64	conv2d_5[0][0]
activation_4 (Activation)	(None, 224, 224, 16)	0	batch_normalization_4[0][0]
conv2d_6 (Conv2D)	(None, 224, 224, 16)	2320	activation_4[0][0]
batch_normalization_5 (BatchNor	(None, 224, 224, 16)	64	conv2d_6[0][0]
activation_5 (Activation)	(None, 224, 224, 16)	0	batch_normalization_5[0][0]
conv2d_7 (Conv2D)	(None, 224, 224, 64)	1088	activation_5[0][0]
add_1 (Add)	(None, 224, 224, 64)	0	add[0][0] conv2d_7[0][0]
batch_normalization_6 (BatchNor	(None, 224, 224, 64)	256	add_1[0][0]
activation_6 (Activation)	(None, 224, 224, 64)	0	batch_normalization_6[0][0]
conv2d_8 (Conv2D)	(None, 224, 224, 16)	1040	activation_6[0][0]
batch_normalization_7 (BatchNor	(None, 224, 224, 16)	64	conv2d_8[0][0]
activation_7 (Activation)	(None, 224, 224, 16)	0	batch_normalization_7[0][0]

activation_7 (Activation)	(None, 224, 224, 16) 0	batch_normalization_7[0][0]
conv2d_9 (Conv2D)	(None, 224, 224, 16) 2320	activation_7[0][0]
batch_normalization_8 (BatchNor	(None, 224, 224, 16) 64	conv2d_9[0][0]
activation_8 (Activation)	(None, 224, 224, 16) 0	batch_normalization_8[0][0]
conv2d_10 (Conv2D)	(None, 224, 224, 64) 1088	activation_8[0][0]
add_2 (Add)	(None, 224, 224, 64) 0	add_1[0][0] conv2d_10[0][0]
batch_normalization_9 (BatchNor	(None, 224, 224, 64) 256	add_2[0][0]
activation_9 (Activation)	(None, 224, 224, 64) 0	batch_normalization_9[0][0]
conv2d_11 (Conv2D)	(None, 112, 112, 64) 4160	activation_9[0][0]
batch_normalization_10 (BatchNo	(None, 112, 112, 64) 256	conv2d_11[0][0]
activation_10 (Activation)	(None, 112, 112, 64) 0	batch_normalization_10[0][0]
conv2d_12 (Conv2D)	(None, 112, 112, 64) 36928	activation_10[0][0]
batch_normalization_11 (BatchNo	(None, 112, 112, 64) 256	conv2d_12[0][0]
activation_11 (Activation)	(None, 112, 112, 64) 0	batch_normalization_11[0][0]
conv2d_14 (Conv2D)	(None, 112, 112, 128) 8320	add_2[0][0]
conv2d_13 (Conv2D)	(None, 112, 112, 128) 8320	activation_11[0][0]
add_3 (Add)	(None, 112, 112, 128) 0	conv2d_14[0][0] conv2d_13[0][0]
batch_normalization_12 (BatchNo	(None, 112, 112, 128) 512	add_3[0][0]
activation_12 (Activation)	(None, 112, 112, 128) 0	batch_normalization_12[0][0]
conv2d_15 (Conv2D)	(None, 112, 112, 64) 8256	activation_12[0][0]
batch_normalization_13 (BatchNo	(None, 112, 112, 64) 256	conv2d_15[0][0]
activation_13 (Activation)	(None, 112, 112, 64) 0	batch_normalization_13[0][0]
conv2d_16 (Conv2D)	(None, 112, 112, 64) 36928	activation_13[0][0]
batch_normalization_14 (BatchNo	(None, 112, 112, 64) 256	conv2d_16[0][0]
activation_14 (Activation)	(None, 112, 112, 64) 0	batch_normalization_14[0][0]
conv2d_17 (Conv2D)	(None, 112, 112, 128) 8320	activation_14[0][0]
add_4 (Add)	(None, 112, 112, 128) 0	add_3[0][0] conv2d_17[0][0]
batch_normalization_15 (BatchNo	(None, 112, 112, 128) 512	add_4[0][0]
activation_15 (Activation)	(None, 112, 112, 128) 0	batch_normalization_15[0][0]
conv2d_18 (Conv2D)	(None, 112, 112, 64) 8256	activation_15[0][0]
batch_normalization_16 (BatchNo	(None, 112, 112, 64) 256	conv2d_18[0][0]

activation_16 (Activation)	(None, 112, 112, 64) 0	batch_normalization_16[0][0]
conv2d_19 (Conv2D)	(None, 112, 112, 64) 36928	activation_16[0][0]
batch_normalization_17 (Batch Normalization)	(None, 112, 112, 64) 256	conv2d_19[0][0]
activation_17 (Activation)	(None, 112, 112, 64) 0	batch_normalization_17[0][0]
conv2d_20 (Conv2D)	(None, 112, 112, 128) 8320	activation_17[0][0]
add_5 (Add)	(None, 112, 112, 128) 0	add_4[0][0] conv2d_20[0][0]
batch_normalization_18 (Batch Normalization)	(None, 112, 112, 128) 512	add_5[0][0]
activation_18 (Activation)	(None, 112, 112, 128) 0	batch_normalization_18[0][0]
conv2d_21 (Conv2D)	(None, 56, 56, 128) 16512	activation_18[0][0]
batch_normalization_19 (Batch Normalization)	(None, 56, 56, 128) 512	conv2d_21[0][0]
activation_19 (Activation)	(None, 56, 56, 128) 0	batch_normalization_19[0][0]
conv2d_22 (Conv2D)	(None, 56, 56, 128) 147584	activation_19[0][0]
batch_normalization_20 (Batch Normalization)	(None, 56, 56, 128) 512	conv2d_22[0][0]
activation_20 (Activation)	(None, 56, 56, 128) 0	batch_normalization_20[0][0]
conv2d_24 (Conv2D)	(None, 56, 56, 256) 33024	add_5[0][0]
conv2d_23 (Conv2D)	(None, 56, 56, 256) 33024	activation_20[0][0]
add_6 (Add)	(None, 56, 56, 256) 0	conv2d_24[0][0] conv2d_23[0][0]
batch_normalization_21 (Batch Normalization)	(None, 56, 56, 256) 1024	add_6[0][0]
activation_21 (Activation)	(None, 56, 56, 256) 0	batch_normalization_21[0][0]
conv2d_25 (Conv2D)	(None, 56, 56, 128) 32896	activation_21[0][0]
batch_normalization_22 (Batch Normalization)	(None, 56, 56, 128) 512	conv2d_25[0][0]
activation_22 (Activation)	(None, 56, 56, 128) 0	batch_normalization_22[0][0]
conv2d_26 (Conv2D)	(None, 56, 56, 128) 147584	activation_22[0][0]
batch_normalization_23 (Batch Normalization)	(None, 56, 56, 128) 512	conv2d_26[0][0]
activation_23 (Activation)	(None, 56, 56, 128) 0	batch_normalization_23[0][0]
conv2d_27 (Conv2D)	(None, 56, 56, 256) 33024	activation_23[0][0]
add_7 (Add)	(None, 56, 56, 256) 0	add_6[0][0] conv2d_27[0][0]
batch_normalization_24 (Batch Normalization)	(None, 56, 56, 256) 1024	add_7[0][0]
activation_24 (Activation)	(None, 56, 56, 256) 0	batch_normalization_24[0][0]
conv2d_28 (Conv2D)	(None, 56, 56, 128) 32896	activation_24[0][0]
batch_normalization_25 (Batch Normalization)	(None, 56, 56, 128) 512	conv2d_28[0][0]

batch_normalization_25 (BatchNormalizatio	(None, 56, 56, 128)	512	conv2d_29[0][0]
activation_25 (Activation)	(None, 56, 56, 128)	0	batch_normalization_25[0][0]
conv2d_29 (Conv2D)	(None, 56, 56, 128)	147584	activation_25[0][0]
batch_normalization_26 (BatchNormalizatio	(None, 56, 56, 128)	512	conv2d_29[0][0]
activation_26 (Activation)	(None, 56, 56, 128)	0	batch_normalization_26[0][0]
conv2d_30 (Conv2D)	(None, 56, 56, 256)	33024	activation_26[0][0]
add_8 (Add)	(None, 56, 56, 256)	0	add_7[0][0] conv2d_30[0][0]
batch_normalization_27 (BatchNormalizatio	(None, 56, 56, 256)	1024	add_8[0][0]
activation_27 (Activation)	(None, 56, 56, 256)	0	batch_normalization_27[0][0]
average_pooling2d (AveragePooling2D)	(None, 7, 7, 256)	0	activation_27[0][0]
flatten (Flatten)	(None, 12544)	0	average_pooling2d[0][0]
dense (Dense)	(None, 2)	25090	flatten[0][0]
=====			
Total params: 871,522			
Trainable params: 866,306			
Non-trainable params: 5,216			

Train on 7303 samples, validate on 2435 samples

Epoch 1/100

7296/7303 [=====>.] - ETA: 0s - loss: 1.1425 - acc: 0.7640

Epoch 00001: val_acc improved from -inf to 0.79630, saving model to drive/My Drive/Colab Notebooks,

7303/7303 [=====] - 129s 18ms/sample - loss: 1.1423 - acc: 0.7641 - val_loss: 0.84846

Epoch 2/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.9031 - acc: 0.8147

Epoch 00002: val_acc improved from 0.79630 to 0.84846, saving model to drive/My Drive/Colab Notebooks,

7303/7303 [=====] - 113s 16ms/sample - loss: 0.9031 - acc: 0.8145 - val_loss: 0.84846

Epoch 3/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.7761 - acc: 0.8292

Epoch 00003: val_acc did not improve from 0.84846

7303/7303 [=====] - 111s 15ms/sample - loss: 0.7760 - acc: 0.8291 - val_loss: 0.84846

Epoch 4/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.6883 - acc: 0.8439

Epoch 00004: val_acc did not improve from 0.84846

7303/7303 [=====] - 110s 15ms/sample - loss: 0.6884 - acc: 0.8438 - val_loss: 0.84846

Epoch 5/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.6274 - acc: 0.8444

Epoch 00005: val_acc did not improve from 0.84846

7303/7303 [=====] - 110s 15ms/sample - loss: 0.6276 - acc: 0.8442 - val_loss: 0.84846

Epoch 6/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.5856 - acc: 0.8513

Epoch 00006: val_acc did not improve from 0.84846

7303/7303 [=====] - 110s 15ms/sample - loss: 0.5854 - acc: 0.8514 - val_loss: 0.84846

Epoch 7/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.5467 - acc: 0.8542

Epoch 00007: val_acc did not improve from 0.84846

7303/7303 [=====] - 111s 15ms/sample - loss: 0.5465 - acc: 0.8543 - val_loss: 0.84846

Epoch 8/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.5235 - acc: 0.8575

Epoch 00008: val_acc improved from 0.84846 to 0.86283, saving model to drive/My Drive/Colab Notebooks,

7303/7303 [=====] - 112s 15ms/sample - loss: 0.5241 - acc: 0.8572 - val_loss: 0.84846

Epoch 9/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.4996 - acc: 0.8576

Epoch 00009: val_acc did not improve from 0.86283

```

Epoch 0000: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4994 - acc: 0.8577 - val_loss: 0.86283
Epoch 10/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4829 - acc: 0.8603
Epoch 00010: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4835 - acc: 0.8602 - val_loss: 0.86283
Epoch 11/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4674 - acc: 0.8643
Epoch 00011: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4676 - acc: 0.8643 - val_loss: 0.86283
Epoch 12/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4496 - acc: 0.8679
Epoch 00012: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4499 - acc: 0.8679 - val_loss: 0.86283
Epoch 13/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4374 - acc: 0.8703
Epoch 00013: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4375 - acc: 0.8703 - val_loss: 0.86283
Epoch 14/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4283 - acc: 0.8688
Epoch 00014: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4284 - acc: 0.8688 - val_loss: 0.86283
Epoch 15/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4060 - acc: 0.8783
Epoch 00015: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.4062 - acc: 0.8781 - val_loss: 0.86283
Epoch 16/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.3908 - acc: 0.8821
Epoch 00016: val_acc did not improve from 0.86283
7303/7303 [=====] - 110s 15ms/sample - loss: 0.3909 - acc: 0.8821 - val_loss: 0.86283
Epoch 17/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.3745 - acc: 0.8912
Epoch 00017: val_acc did not improve from 0.86283
7303/7303 [=====] - 111s 15ms/sample - loss: 0.3751 - acc: 0.8909 - val_loss: 0.86283
Epoch 18/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.3621 - acc: 0.8957
Epoch 00018: val_acc did not improve from 0.86283
7303/7303 [=====] - 111s 15ms/sample - loss: 0.3623 - acc: 0.8955 - val_loss: 0.86283
Epoch 00018: early stopping
Create Model Resnet 2022.3518 seconds

```

```
watch_restart()
```

```
# ===== 4. Evaluate Models
```

```

best_model_resnet = get_model_resnet(input_shape=IMAGE_INPUT_SHAPE, depth=NUM_DEPTH, num_classes=NUM_CLASSES)
best_model_resnet.load_weights(MODEL_NAME_RESNET)
best_model_resnet.compile(loss='categorical_crossentropy',
                          optimizer=Adam(lr=lr_schedule(NUM_EPOCHS)),
                          metrics=['accuracy'])

```

```

def print_score(title, model, X_test, y_test, label_class):
    scores = model.evaluate(X_test, y_test, verbose=1)
    print("{0} {1} {2}".format(title, model.metrics_names[1], scores[1]*100))
    y_pred = model.predict(X_test)

```

```
print('\n', classification_report(np.where(y_test > 0)[1],
                                   np.argmax(y_pred, axis=1),
                                   target_names=list(label_class.keys())), sep='')

def plot_accuracy(title, history):
    plt.figure(figsize=(8,8))
    plt.plot(history.history['acc'])      # Training Accuracy
    plt.plot(history.history['val_acc']) # Validation Accuracy
    plt.title('{0} Model Accuracy'.format(title))
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()

def plot_loss(title, history):
    plt.figure(figsize=(8,8))
    plt.plot(history.history['loss'])      # Training Loss
    plt.plot(history.history['val_loss']) # Validation Loss
    plt.title('{0} Model Loss'.format(title))
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()

# Scores
print_score('Resnet', best_model_resnet, X_test, y_test, LABEL_CLASS)

# Plot Accuracy
plot_accuracy('Resnet', history_resnet)

# Plot Model Loss
plot_loss('Resnet', history_resnet)

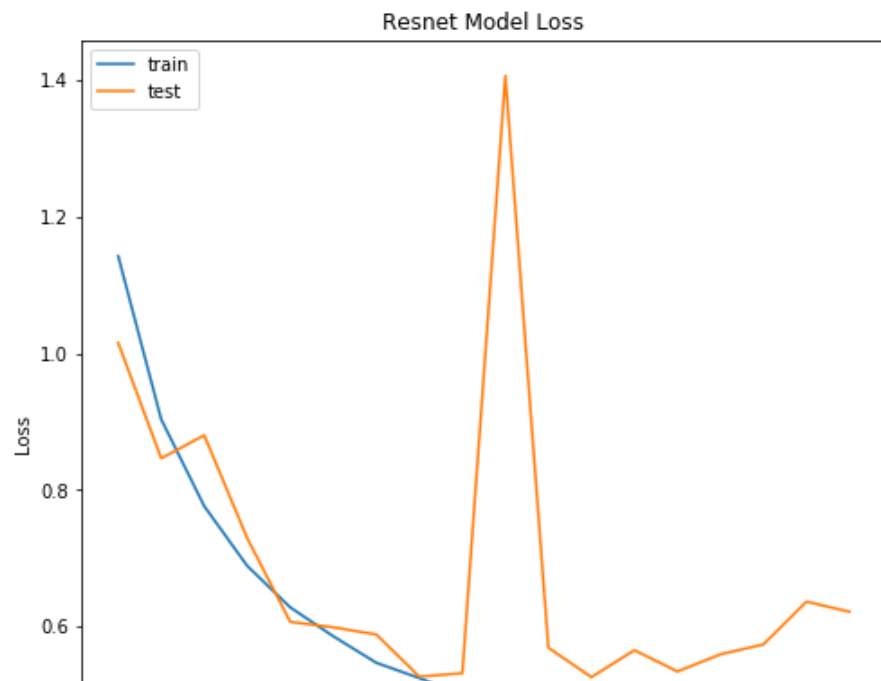

watch_print('Evaluate Model')
```




2435/2435 [=====] - 11s 5ms/sample - loss: 0.5262 - acc: 0.8628

Resnet acc 86.28336787223816

	precision	recall	f1-score	support
Down	0.88	0.79	0.83	1046
Up	0.85	0.92	0.88	1389
accuracy			0.86	2435
macro avg	0.87	0.85	0.86	2435
weighted avg	0.86	0.86	0.86	2435





Evaluate Model 34.9093 seconds