

```
# ===== 0.1 Check for GPU
!nvidia-smi

import tensorflow as tf
from tensorflow.python.client import device_lib

print("Tensorflow Version:", tf.__version__)
print("Is Build with CUDA:", tf.test.is_built_with_cuda())
print("Is GPU available:", tf.test.is_gpu_available(cuda_only=False, min_cuda_compute_capability=None))
print(device_lib.list_local_devices())
```



Saved successfully!



+-----+-----+-----+									
NVIDIA-SMI 440.48.02		Driver Version: 418.67				CUDA Version: 10.1			
+-----+-----+-----+									
GPU Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC	
Fan Temp Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util		Compute M.	
=====+=====+=====									
0 Tesla P100-PCIE...		Off		00000000:00:04.0 Off		0			
N/A 36C P0		26W / 250W		0MiB / 16280MiB		0%		Default	
+-----+-----+-----+									

Processes:				GPU Memory Usage
GPU	PID	Type	Process name	
No running processes found				

We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %tensorflow_vers

research.google.com/drive/1SF-J

research.google.com/drive/1SF-J

```

import pandas as pd
import numpy as np
import cv2
import base64
import matplotlib.pyplot as plt
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.applications import vgg16, vgg19, resnet50
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D, GlobalAveragePooling2D
from tensorflow.keras.layers import BatchNormalization, Activation, Input, AveragePooling2D
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping, LearningRateScheduler, ReduceLROnPlateau
from tensorflow.keras.regularizers import l2

```

```

PATH_ROOT = 'drive/My Drive/Colab Notebooks'
PATH_FILE = PATH_ROOT + '/Datasets/ChartImages/charts_1_day_50_periods.csv'
MODEL_NAME_VGG16 = PATH_ROOT + '/Models/ChartImages/model_vgg16_1_day_50_periods.h5'
MAX_RGB = 255
IMAGE_DIMENSION = 224
IMAGE_INPUT_SHAPE = (IMAGE_DIMENSION, IMAGE_DIMENSION, 3)
LABEL_CLASS = { 'Down': 0, 'Up': 1 }
NUM_CLASSES = 2
NUM_BATCH_SIZE_VGG16 = 128
NUM_EPOCHS = 100

```

Utils

Saved successfully!



```

global start_time
start_time = time.time()
def watch_print(title):
    global start_time
    print(title, round(time.time() - start_time, 4), 'seconds')
def memory_print():
    memory = dict(psutil.virtual_memory()._asdict())
    print("Memory Capacity", memory['total'] >> 30, "GB")
    print("Memory Left", memory['free'] >> 30, "GB")
    print("Memory Used", memory['used'] >> 30, "GB")
    print("Memory Used:", memory['percent'], "percent")

```

Callbacks for model saving and stopping.

Training should be stopped when val_acc (validation accuracy) clearly stops increasing to prevent overfitting

```

checkpoint_vgg16 = ModelCheckpoint(filepath=MODEL_NAME_VGG16, monitor='val_acc', verbose=1, save_best_only=True)
early_stopping = EarlyStopping(monitor='val_acc', patience=10, verbose=1)

```

watch_restart()

===== 1.0 Get Image Data from CSV

```

df = pd.read_csv(PATH_FILE)
print(df.head())

```

```

memory_print()
watch_print('Get Data')

```



	Id	Symbol	...	Y_Prediction	X_Image
0	1	M	...	1	iVBORw0KGgoAAAANSU...EugAAAOAAAADgCAYAAAAaLWrhAA...
1	2	M	...	1	iVBORw0KGgoAAAANSU...EugAAAOAAAADgCAYAAAAaLWrhAA...
2	3	M	...	1	iVBORw0KGgoAAAANSU...EugAAAOAAAADgCAYAAAAaLWrhAA...
3	4	M	...	1	iVBORw0KGgoAAAANSU...EugAAAOAAAADgCAYAAAAaLWrhAA...
4	5	M	...	1	iVBORw0KGgoAAAANSU...EugAAAOAAAADgCAYAAAAaLWrhAA...

```
[5 rows x 5 columns]
Memory Capacity 25 GB
Memory Left 22 GB
Memory Used 1 GB
Memory Used: 4.7 percent
Get Data 4.5841 seconds
```

```
watch_restart()
```

```
# ===== 2.0 Prepare Data
```

```
# Set up X and y
```

```
items = []
```

```
for index, row in df.iterrows():
```

```
    # Convert from base64 string to byte array
```

```
    item_byte_array = base64.b64decode(df['X_Image'][index])
```

```
    # Convert byte array to numpy array for OpenCv usage
```

```
    item_np = np.frombuffer(item_byte_array, dtype=np.uint8)
```

Saved successfully!



```
image
np, flags=1)
```

```
items.append(item_image)
```

```
# if index < 1:
```

```
    # plt.imshow(cv2.cvtColor(item_image, cv2.COLOR_BGR2RGB));
```

```
    # plt.show()
```

```
X = np.array(items)
```

```
y = to_categorical(df[['Y_Prediction']].values)
```

```
df = None # Clear RAM
```

```
print('y type', type(y))
```

```
print('y shape', y.shape)
```

```
print('X type', type(X))
```

```
print('X shape', X.shape)
```

```
# Normalize input data.
```

```
# Neural Networks work best when input data are between 0 and 1 (Instead of 0 to 255).
```

```
X = X / MAX_RGB
```

```
# Split Train and Test
```

```
def split(X, y, proportion):
```

```
    ratio = int(X.shape[0]/proportion)
```

```
    X_train = X[ratio:,:]
```

```
    X_test = X[:ratio,:]
```

```
    y_train = y[ratio:,:]
```

```
    y_test = y[:ratio,:]
```

```
    return X_train, X_test, y_train, y_test
```

```
# X train. X test. y train. y test = split(X, y, 4) # Uses less RAM
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42) # Uses a lot of RAM
X = y = None # Clear RAM

```

```

print("X_train Shape", X_train.shape)
print("y_train Shape", y_train.shape)
print("X_test Shape", X_test.shape)
print("y_test Shape", y_test.shape)

```

```

memory_print()
watch_print('Prepare Data')

```

```

y type <class 'numpy.ndarray'>
y shape (9738, 2)
X type <class 'numpy.ndarray'>
X shape (9738, 224, 224, 3)
X_train Shape (7303, 224, 224, 3)
y_train Shape (7303, 2)
X_test Shape (2435, 224, 224, 3)
y_test Shape (2435, 2)
Memory Capacity 25 GB
Memory Left 0 GB
Memory Used 24 GB
Memory Used: 52.2 percent
Prepare Data 17.8596 seconds

```

```
watch_restart()
```

```
# ----- 3.0 Create Model VGG16
```

```
def create_model_vgg16(input_shape, num_classes):
```

```

    # Remove top since different number of output classes than the pretrained model.
    model_pretrained = vgg16.VGG16(include_top=False, input_shape=input_shape)

```

```

    # Freeze layers since they have already been pretrained.
    for layer in model_pretrained.layers:
        layer.trainable = False
    ret.add(layer)

```

```

    # Flatten
    ret.add(Flatten())

```

```

    ret.add(Conv2D(512, kernel_size = (3,3), padding = 'valid'))
    ret.add(GlobalAveragePooling2D())

```

```

    # 4 classes
    ret.add(Dense(num_classes, activation='softmax'))

```

```
    return ret
```

```

opt = SGD(lr=0.001)
model_vgg16 = get_model_vgg16(IMAGE_INPUT_SHAPE, NUM_CLASSES)
model_vgg16.compile(loss='categorical_crossentropy',
                    #optimizer=opt,
                    optimizer='adam',
                    metrics=['accuracy'])
model_vgg16.summary()

```

```
checkpoint_vgg16 = ModelCheckpoint(filepath=MODEL_NAME_VGG16, monitor='val_acc', verbose=1, save_best_only=True)
history_vgg16 = model_vgg16.fit(X_train,
                                y_train,
                                batch_size=NUM_BATCH_SIZE_VGG16,
                                epochs=NUM_EPOCHS,
                                validation_data=(X_test, y_test),
                                shuffle=True,
                                callbacks=[checkpoint_vgg16, early_stopping])

memory_print()
watch_print('Create Model VGG16')
```



Saved successfully!



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python/ops/resource_instructions.py:1660: tf.nn.conv2d is deprecated and will be removed in a future version. Instructions for updating:
 If using Keras pass *_constraint arguments to layers.
 Downloading data from https://github.com/fchollet/deep-learning-models/releases/download/v0.1/vgg16_weights_tf_dim_ordering_tf_data_format.h5
 58892288/58889256 [=====] - 5s 0us/step
 Model: "sequential"

Layer (type)	Output Shape	Param #
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 5, 5, 512)	2359808
global_average_pooling2d (GlobalAveragePooling2D)	(None, 512)	0
dense (Dense)	(None, 2)	1026
Total params: 17,075,522		
Trainable params: 2,360,834		
Non-trainable params: 14,714,688		

Saved successfully!



Train on 7303 samples, validate on 2435 samples

Epoch 1/100

7296/7303 [=====>.] - ETA: 0s - loss: 2.0440 - acc: 0.5525

Epoch 00001: val_acc improved from -inf to 0.68665, saving model to drive/My Drive/Colab Notebooks

7303/7303 [=====] - 42s 6ms/sample - loss: 2.0428 - acc: 0.5525 - val_loss: 0.68665

Epoch 2/100

7296/7303 [=====>.] - ETA: 0s - loss: 0.6038 - acc: 0.6859

Epoch 00002: val_acc did not improve from 0.68665

7303/7303 [=====] - 78s 4ms/sample - loss: 0.6035 - acc: 0.6860 - val_loss: 0.68665

```
7303/7303 [=====] - 28s 4ms/sample - loss: 0.5805 - acc: 0.7076 - val_loss
Epoch 3/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5807 - acc: 0.7076
Epoch 00003: val_acc did not improve from 0.68665
7303/7303 [=====] - 28s 4ms/sample - loss: 0.5805 - acc: 0.7078 - val_loss
Epoch 4/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5660 - acc: 0.7196
Epoch 00004: val_acc did not improve from 0.68665
7303/7303 [=====] - 28s 4ms/sample - loss: 0.5662 - acc: 0.7194 - val_loss
Epoch 5/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5484 - acc: 0.7392
Epoch 00005: val_acc did not improve from 0.68665
7303/7303 [=====] - 28s 4ms/sample - loss: 0.5484 - acc: 0.7391 - val_loss
Epoch 6/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5631 - acc: 0.7348
Epoch 00006: val_acc improved from 0.68665 to 0.72936, saving model to drive/My Drive/Colab Notebook
7303/7303 [=====] - 32s 4ms/sample - loss: 0.5633 - acc: 0.7348 - val_loss
Epoch 7/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5649 - acc: 0.7235
Epoch 00007: val_acc did not improve from 0.72936
7303/7303 [=====] - 28s 4ms/sample - loss: 0.5646 - acc: 0.7237 - val_loss
Epoch 8/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5602 - acc: 0.7288
Epoch 00008: val_acc improved from 0.72936 to 0.78973, saving model to drive/My Drive/Colab Notebook
7303/7303 [=====] - 33s 4ms/sample - loss: 0.5599 - acc: 0.7290 - val_loss
Epoch 9/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5043 - acc: 0.7812
Epoch 00009: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.5043 - acc: 0.7812 - val_loss
Epoch 10/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5000 - acc: 0.7808
Epoch 00010: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.5001 - acc: 0.7808 - val_loss
Epoch 11/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5169 - acc: 0.7701
Epoch 00011: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.5169 - acc: 0.7701 - val_loss
Epoch 12/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4922 - acc: 0.7896
Epoch 00012: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4922 - acc: 0.7897 - val_loss
Epoch 13/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4917 - acc: 0.7876
Epoch 00013: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4916 - acc: 0.7876 - val_loss
Epoch 14/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5136 - acc: 0.7681
Epoch 00014: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.5138 - acc: 0.7680 - val_loss
Epoch 15/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4758 - acc: 0.8024
Epoch 00015: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4757 - acc: 0.8023 - val_loss
Epoch 16/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5203 - acc: 0.7566
Epoch 00016: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.5203 - acc: 0.7567 - val_loss
Epoch 17/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4743 - acc: 0.8036
Epoch 00017: val_acc did not improve from 0.78973
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4742 - acc: 0.8035 - val_loss
Epoch 18/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4586 - acc: 0.8113
Epoch 00018: val_acc improved from 0.78973 to 0.79548, saving model to drive/My Drive/Colab Notebook
7303/7303 [=====] - 32s 4ms/sample - loss: 0.4587 - acc: 0.8113 - val_loss
```

Saved successfully!




```

7303/7303 [=====] - 32s 4ms/sample - loss: 0.5045 - acc: 0.7749 - val_loss
Epoch 19/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5043 - acc: 0.7749
Epoch 00019: val_acc improved from 0.79548 to 0.80739, saving model to drive/My Drive/Colab Notebo
7303/7303 [=====] - 32s 4ms/sample - loss: 0.5045 - acc: 0.7749 - val_loss
Epoch 20/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4674 - acc: 0.8076
Epoch 00020: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4676 - acc: 0.8075 - val_loss
Epoch 21/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4538 - acc: 0.8141
Epoch 00021: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4538 - acc: 0.8142 - val_loss
Epoch 22/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.5011 - acc: 0.7726
Epoch 00022: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.5010 - acc: 0.7727 - val_loss
Epoch 23/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4680 - acc: 0.8035
Epoch 00023: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4681 - acc: 0.8032 - val_loss
Epoch 24/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4619 - acc: 0.8050
Epoch 00024: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4617 - acc: 0.8051 - val_loss
Epoch 25/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4526 - acc: 0.8120
Epoch 00025: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4524 - acc: 0.8120 - val_loss
Epoch 26/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4571 - acc: 0.8072
Epoch 00026: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4571 - acc: 0.8071 - val_loss
Epoch 27/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4829 - acc: 0.7917
Epoch 00027: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4831 - acc: 0.7913 - val_loss
Epoch 28/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4669 - acc: 0.8007
Epoch 00028: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4668 - acc: 0.8008 - val_loss
Epoch 29/100
7296/7303 [=====>.] - ETA: 0s - loss: 0.4658 - acc: 0.8000
Epoch 00029: val_acc did not improve from 0.80739
7303/7303 [=====] - 27s 4ms/sample - loss: 0.4657 - acc: 0.8001 - val_loss
Epoch 00029: early stopping
Memory Capacity 25 GB
Memory Left 0 GB
Memory Used 22 GB
Memory Used: 59.4 percent
Create Model VGG16 835.5362 seconds

```

Saved successfully!



```
watch_restart()
```

```
# ===== 4. Evaluate Models
```

```
best_model_vgg16 = get_model_vgg16(IMAGE_INPUT_SHAPE, NUM_CLASSES)
```

```
best_model_vgg16.load_weights(MODEL_NAME_VGG16)
```

```
best_model_vgg16.compile(loss='categorical_crossentropy',
```

```

best_model_vgg16.compile(loss=categorical_crossentropy,
                        # optimizer=opt,
                        optimizer='adam',
                        metrics=['accuracy'])

def print_score(title, model, X_test, y_test, label_class):
    scores = model.evaluate(X_test, y_test, verbose=1)
    print("{0} {1} {2}".format(title, model.metrics_names[1], scores[1]*100))
    y_pred = model.predict(X_test)
    print('\n', classification_report(np.where(y_test > 0)[1],
                                        np.argmax(y_pred, axis=1),
                                        target_names=list(label_class.keys())), sep='')

def plot_accuracy(title, history):
    plt.figure(figsize=(8,8))
    plt.plot(history.history['acc'])      # Training Accuracy
    plt.plot(history.history['val_acc']) # Validation Accuracy
    plt.title('{0} Model Accuracy'.format(title))
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['train', 'test'], loc='upper left')
    plt.show()

def plot_loss(title, history):
    plt.figure(figsize=(8,8))
    plt.plot(history.history['loss'])      # Training Loss
    plt.plot(history.history['val_loss']) # Validation Loss
    plt.title('{0} Model Loss'.format(title))

    plt.legend(['train', 'test'], loc='upper left')
    plt.show()

# Scores
print_score('VGG16', best_model_vgg16, X_test, y_test, LABEL_CLASS)

# Plot Accuracy
plot_accuracy('VGG16', history_vgg16)

# Plot Model Loss
plot_loss('VGG16', history_vgg16)

watch_print('Evaluate Model')


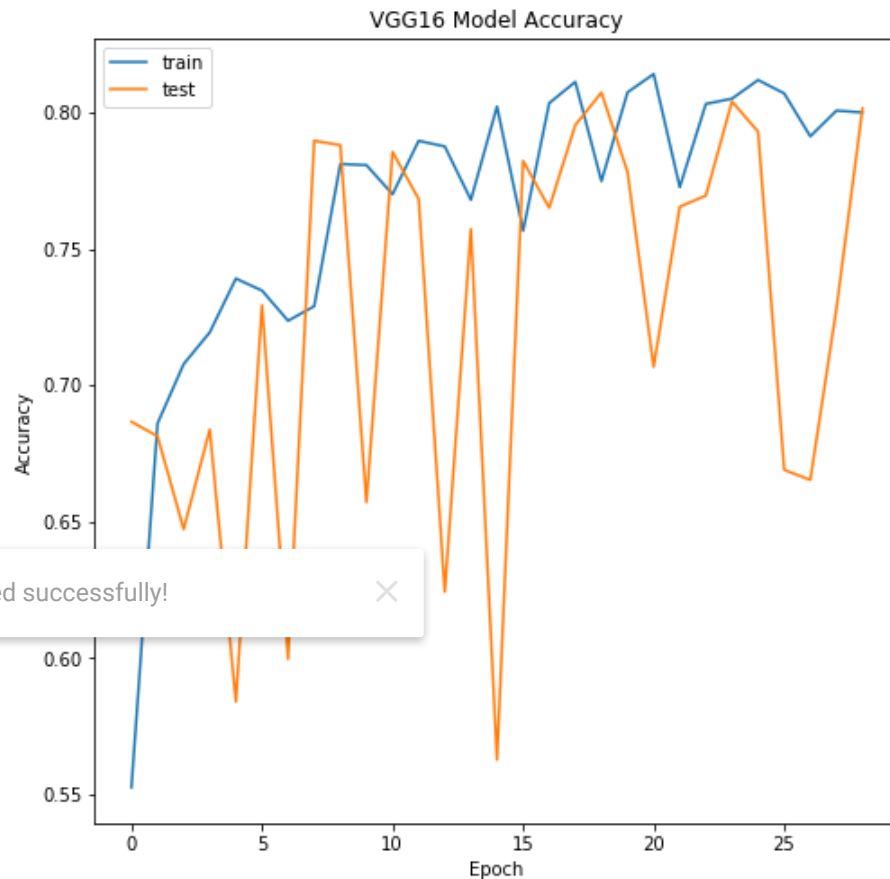
```


Saved successfully!



2435/2435 [=====] - 8s 3ms/sample - loss: 0.4888 - acc: 0.8074
VGG16 acc 80.73921799659729

	precision	recall	f1-score	support
Down	0.82	0.71	0.76	1046
Up	0.80	0.88	0.84	1389
accuracy			0.81	2435
macro avg	0.81	0.79	0.80	2435
weighted avg	0.81	0.81	0.80	2435





Evaluate Model 36.243 seconds

```
# !pip freeze > "drive/My Drive/Colab Notebooks/requirements_ChartsPrediction_VGG16_1_day_50_periods.tx"
```

Saved successfully!

